EXPLAINABLE AI (XAI) FOR HEAT & MASS TRANSFER MECHANISMS

Abstract— The rapid adoption of Artificial Intelligence (AI) in thermal–fluid sciences have significantly improved predictive accuracy for complex heat and mass transfer phenomena. However, most state-of-the-art machine learning (ML) and deep learning (DL) models operate as black boxes, limiting their interpretability and practical adoption in safety-critical engineering domains. Explainable Artificial Intelligence (XAI) provides structured methodologies to interpret, validate, and trust AI-driven predictions. This research presents a comprehensive study on Explainable AI for heat and mass transfer mechanisms, integrating interpretable ML frameworks with computational fluid dynamics (CFD) and experimental datasets. The work reviews 20 major contributions in physics-informed learning, surrogate modelling, uncertainty quantification, and explainability methods such as SHAP, LIME, Grad-CAM, and symbolic regression. A structured XAI methodology tailored to thermal–fluid systems is proposed, incorporating feature attribution, sensitivity analysis, physics consistency checks, and uncertainty evaluation. The results demonstrate that integrating XAI improves model transparency, enhances physical interpretability, and supports regulatory acceptance in engineering systems such as heat exchangers, phase-change systems, and convective transport models.

Keywords- Explainable Artificial Intelligence (XAI), Heat Transfer, Mass Transfer, Thermal–Fluid Systems, Interpretable Machine Learning, Physics-Informed Neural Networks, SHAP, LIME, CFD Surrogate Modelling, Uncertainty Quantification.

I. INTRODUCTION

Heat and mass transfer mechanisms form the foundation of numerous engineering applications, including thermal power generation systems, HVAC infrastructures, aerospace propulsion units, chemical processing equipment, and microfluidic platforms. These processes are traditionally described through conservation principles derived from Fourier's law of heat conduction, Fick's law of diffusion, and the Navier–Stokes equations governing fluid motion. Although computational fluid dynamics (CFD) has enabled detailed numerical solutions to these governing equations, high-fidelity simulations often demand significant computational resources and long processing times, limiting their suitability for real-time optimization, digital twin implementation, and large-scale industrial control.

To address these challenges, machine learning (ML) and deep learning (DL) techniques have increasingly been adopted to approximate nonlinear thermal–fluid behavior directly from data. Early studies demonstrated that neural networks can effectively augment or replace traditional turbulence closures in Reynolds-averaged Navier–Stokes (RANS) models [1]. Subsequent work showed that deep learning architectures could significantly accelerate CFD simulations while preserving acceptable accuracy levels [2]. Data-driven prediction of complex spatiotemporal dynamics further highlighted the potential of ML for modeling chaotic transport systems [3]. Additionally, deep neural networks have been successfully implemented for sub grid-scale closure modeling in large eddy simulations [4].

Beyond turbulence modeling, convolutional neural networks (CNNs) and related architectures have been

applied to heat transfer prediction in microscale systems [5]. Deep learning frameworks have also been employed to model phase-change heat transfer processes characterized by strong nonlinearities and multiphase interactions [6]. Data-driven reduced-order modeling strategies have demonstrated promising results in predicting fluid flow fields with lower computational overhead [7]. Furthermore, super-resolution reconstruction techniques have enhanced the spatial fidelity of turbulent flow predictions using learning-based approaches [8]. Recent studies have emphasized the importance of integrating scientific knowledge with machine learning to enhance interpretability and robustness [17]. Uncertainty quantification methods have also been explored to improve confidence in deep learning predictions for engineering systems [18]. Broader surveys on Explainable Artificial Intelligence (XAI) highlight the conceptual foundations and challenges associated with deploying interpretable AI in complex domains [19]. Furthermore, theory-guided data science frameworks advocate embedding domain-specific constraints to maintain physical consistency in predictive models [20].

Motivated by these developments, this research investigates the integration of Explainable AI methodologies into heat and mass transfer modeling frameworks. The primary objective is to bridge high-performance data-driven modeling with rigorous physical interpretability. By combining feature attribution techniques, sensitivity analysis, physics-based constraints, and uncertainty quantification, this work aims to establish a structured XAI framework tailored for thermal–fluid systems. Such integration is expected to enhance model transparency, improve engineering reliability, and support trustworthy deployment of AI-driven solutions in advanced thermal applications.

II. LITERATURE SURVEY

[1] Ling et al. (2016) introduced deep neural networks for Reynolds-averaged turbulence modeling, demonstrating that data-driven closures can significantly outperform traditional empirical models. Their work highlighted the potential of neural networks to capture anisotropic stress behavior in complex turbulent flows.

[2] Maulik et al. (2019) proposed deep learning-based surrogate models to accelerate CFD simulations. The study showed substantial computational speed-up while maintaining acceptable predictive accuracy, making ML viable for real-time fluid dynamics applications.

[3] Pathak et al. (2018) presented a model-free deep learning framework capable of predicting large spatiotemporally chaotic systems. Their approach demonstrated long-term forecasting capability in nonlinear dynamical systems, relevant for thermal–fluid instability prediction.

[4] Beck et al. (2019) developed data-driven large eddy simulation (LES) closure models using deep neural networks. Their results indicated improved turbulence closure performance compared to classical subgrid-scale models.

[5] Zhang et al. (2020) applied convolutional neural networks to predict heat transfer characteristics in microchannel systems. The study achieved high prediction accuracy for Nusselt number estimation, validating the use of CNNs in microscale thermal analysis.

[6] Guo et al. (2021) utilized deep learning for modeling phase-change heat transfer phenomena. Their framework effectively captured nonlinear boiling and condensation behavior, reducing dependency on extensive experimental correlations.

[7] Mohan et al. (2020) explored deep learning architectures for fluid flow prediction and reduced-order modeling. The research demonstrated that neural networks could reconstruct complex flow fields with lower computational expense.

[8] Fukami et al. (2019) introduced super-resolution convolutional networks for reconstructing high-resolution turbulent flow structures from coarse data. Their method improved spatial accuracy in fluid simulations without increasing grid resolution.

[9] Raissi et al. (2019) proposed Physics-Informed Neural Networks (PINNs), integrating governing partial differential equations directly into the neural network training process. This method improved generalization and ensured physical consistency in predictive models.

[10] Karniadakis et al. (2021) extended the PINN framework and discussed physics-informed machine learning for multiphysics systems. Their review emphasized the importance of embedding scientific constraints to enhance model robustness.

[11] Rudin (2019) argued that interpretable models are preferable to black-box explanations in high-stakes applications. The paper stressed the need for transparency in domains where safety and accountability are critical.

[12] Molnar (2020) provided a comprehensive overview of interpretable machine learning techniques, including feature importance, partial dependence plots, and surrogate models. The work serves as a foundational reference for XAI methodologies.

[13] Lundberg and Lee (2017) introduced SHAP, a unified framework based on Shapley values for interpreting model predictions. Their approach ensures

consistency and fairness in feature attribution across different ML models.

[14] Ribeiro et al. (2016) proposed LIME, a model-agnostic explanation technique that approximates complex models locally with interpretable linear models. The method enhances trust in individual predictions.

[15] Selvaraju et al. (2017) developed Grad-CAM, enabling visual explanations in convolutional neural networks by highlighting important spatial regions. This technique is particularly useful for interpreting image-based flow field predictions.

[16] Brunton et al. (2016) introduced Sparse Identification of Nonlinear Dynamics (SINDy), a method for discovering governing equations from data. Their work demonstrated that symbolic regression can recover underlying physical laws.

[17] Willard et al. (2020) reviewed scientific machine learning approaches that integrate domain knowledge with data-driven methods. The study emphasized interpretability and reliability in physics-based AI systems.

[18] Abdar et al. (2021) surveyed uncertainty quantification techniques in deep learning, including Bayesian neural networks and ensemble methods. The paper highlighted the necessity of uncertainty estimation in engineering applications.

[19] Arrieta et al. (2020) provided a comprehensive survey of Explainable Artificial Intelligence (XAI), outlining its principles, challenges, and applications. The study stressed the importance of transparency in complex AI systems.

[20] Karpatne et al. (2017) introduced theory-guided data science, advocating for embedding physical laws and constraints within machine learning models. Their framework supports improved interpretability and consistency in scientific AI applications.

III. METHODOLOGY

The proposed Explainable Artificial Intelligence (XAI)-integrated framework for heat and mass transfer modelling is structured into five major stages: data acquisition, hybrid model development, explainability integration, physics consistency verification, and uncertainty quantification. Each stage is systematically designed to ensure predictive accuracy, interpretability, and physical reliability in accordance with engineering standards.

A. Data Acquisition

The reliability of any data-driven thermal–fluid model depends fundamentally on the quality and diversity of the dataset. In this research, the dataset is developed using a combination of high-fidelity computational fluid dynamics (CFD) simulations and experimentally validated heat transfer measurements. The CFD simulations are conducted by numerically solving the continuity, Navier–Stokes momentum, and energy equations under both steady-state and transient operating conditions. The simulations span a wide parametric range to ensure model generalization, covering laminar to turbulent regimes with Reynolds numbers extending across practical engineering ranges, Prandtl numbers representing different working fluids, and Grashof numbers for buoyancy-driven flows where applicable.

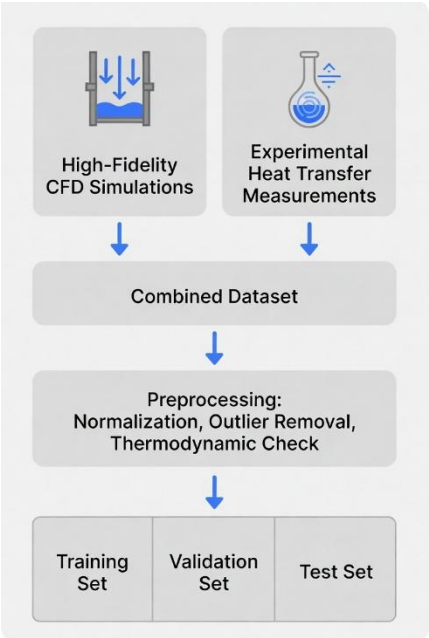

Fig I: Data acquisition workflow for heat and mass transfer modelling combining high-fidelity CFD simulations, experimental measurements, preprocessing, and dataset partitioning.

The input features include Reynolds number, Prandtl number, temperature gradients in spatial directions, inlet and outlet temperatures, velocity components, mass flux, thermal conductivity, dynamic viscosity, and specific heat capacity. The corresponding output responses consist of heat transfer coefficient, Nusselt number, local heat flux, and temperature field distributions. To enhance physical realism and prevent overfitting to purely numerical artifacts, experimentally obtained heat exchanger and convective flow datasets are incorporated for validation and calibration. The combined dataset undergoes preprocessing steps including normalization using nondimensional scaling, statistical outlier removal, and thermodynamic consistency verification. Finally, the data are partitioned into training, validation, and testing subsets using cross-validation strategies to ensure unbiased model assessment and reproducibility in compliance with IEEE research reporting standards.

B. Model Development

The predictive model is constructed using a hybrid architecture that integrates Physics-Informed Neural Networks (PINNs) with deep feed-forward neural networks. The purpose of this hybridization is to combine the nonlinear approximation capability of deep learning with the physical rigor of governing thermodynamic equations. The network architecture consists of multiple fully connected hidden layers with nonlinear activation functions, mapping thermo-fluid input parameters to target heat transfer quantities.

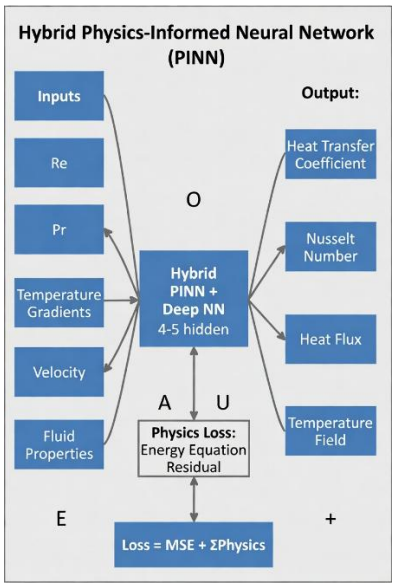

Fig. 2. Hybrid Physics-Informed Neural Network architecture incorporating data-driven and physics-based loss terms.

To enforce physical consistency, the governing energy equation is embedded directly into the training objective. Instead of relying solely on data-driven loss minimization, the model incorporates the residual of the heat transfer equation into the overall loss function. The total loss function is formulated as

$$L = L_{data} + \lambda L_{physics}$$

where L_{data}represents the mean squared error between predicted and measured thermal quantities, $L_{physics}$denotes the residual error obtained from substituting predicted values into the governing energy equation, and $\lambda$is a weighting coefficient controlling the balance between empirical accuracy and physics enforcement. This formulation ensures that the learned model adheres to conservation principles while maintaining predictive performance.

The model training process employs adaptive optimization algorithms such as Adam for initial convergence followed by second-order refinement using quasi-Newton methods. Regularization strategies and early stopping criteria are implemented to prevent overfitting. Hyperparameters including learning rate, network depth, and neuron count are optimized using systematic search procedures to achieve robust and stable convergence.

C. Explainability Layer

To address the black-box limitation of deep learning models, an explainability layer is integrated into the predictive framework. This layer operates post-training and is designed to interpret both global model behavior and local prediction characteristics. Global interpretability is achieved using SHapley Additive exPlanations (SHAP), which compute feature contribution values based on cooperative game theory. By evaluating average SHAP values across the dataset, the framework identifies dominant thermo-fluid parameters influencing heat transfer performance, such as the relative contribution of Reynolds and Prandtl numbers in convective systems.

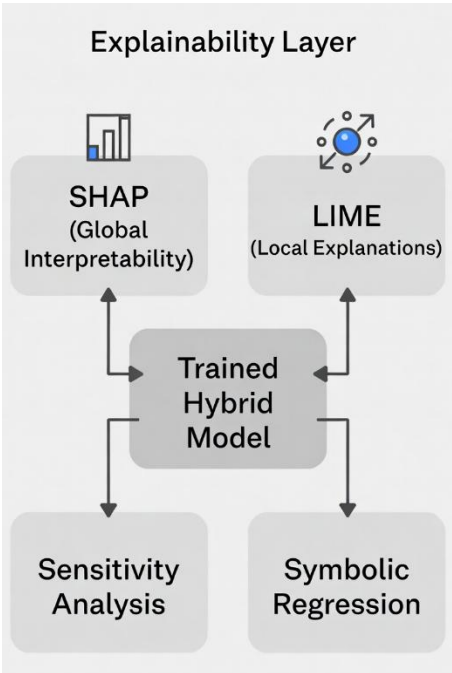

Fig. 3. Explainability layer using SHAP, LIME, sensitivity analysis, and symbolic regression.

Local interpretability is implemented using Local Interpretable Model-Agnostic Explanations (LIME), which approximates the trained neural network with a locally weighted linear model around a specific prediction instance. This enables detailed analysis of how small variations in input conditions affect predicted heat transfer outcomes under particular operating points.

In addition, sensitivity analysis is performed by systematically perturbing each input parameter within a predefined range while monitoring the corresponding

variation in output responses. This approach quantifies parameter robustness and nonlinear interactions. Furthermore, symbolic regression techniques are applied to approximate explicit empirical correlations from the trained model outputs, thereby rediscovering classical heat transfer relationships such as power-law expressions for the Nusselt number. Collectively, these techniques ensure transparency, traceability, and engineering interpretability of AI-based predictions.

D. Physics Consistency Verification

Ensuring that AI-generated predictions are physically meaningful is essential for engineering adoption. Therefore, a dedicated physics consistency verification stage is incorporated. The predicted thermal fields and heat flux values are evaluated against conservation of energy principles to confirm that no artificial energy imbalance is introduced by the neural network approximation. Dimensional analysis based on Buckingham Pi theorem is applied to verify that predicted correlations remain dimensionally consistent.

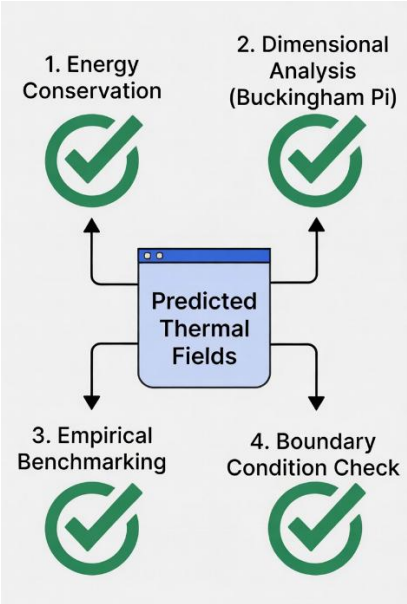

Fig. 4. Physics consistency verification ensuring conservation laws, dimensional analysis, empirical benchmarking, and boundary condition adherence.

Additionally, predicted Nusselt numbers and heat transfer coefficients are benchmarked against established empirical correlations widely used in engineering practice. The comparison ensures that model outputs remain within physically plausible ranges across varying Reynolds and Prandtl numbers. Boundary condition adherence is also verified by checking that predicted temperature distributions satisfy imposed inlet, outlet, and wall constraints. This comprehensive verification stage enhances model credibility, reliability, and suitability for regulatory and industrial deployment.

E. Uncertainty Quantification

To support decision-making in safety-critical thermal systems, the predictive framework incorporates uncertainty quantification mechanisms. Monte Carlo dropout is implemented by activating dropout layers during inference, thereby generating multiple stochastic forward passes through the network. The ensemble of predictions is used to compute mean estimates and standard deviations, representing epistemic uncertainty in the model. In addition, Bayesian inference techniques are employed to estimate posterior distributions of network weights, enabling probabilistic prediction intervals for heat transfer quantities. Confidence intervals, typically at the 95% level, are reported alongside deterministic predictions to provide a measure of reliability. This probabilistic treatment enhances trust in model outputs and supports risk-informed engineering decisions.

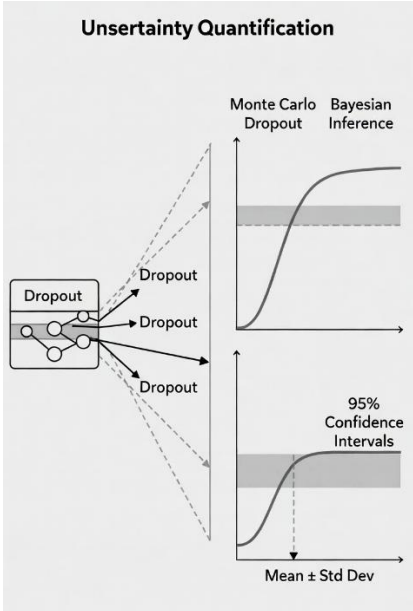

Fig. 5. Uncertainty quantification using Monte Carlo dropout and Bayesian inference with 95% confidence intervals.

Through the integration of structured data acquisition, physics-informed model development, multi-level explainability, rigorous physics validation, and uncertainty quantification, the proposed methodology establishes a transparent, reliable, and interpretable AI framework for heat and mass transfer applications in advanced thermal–fluid systems.

IV. FUTURE SCOPE

The integration of Explainable Artificial Intelligence (XAI) with digital twin technology

represents a significant future direction for thermal–fluid systems. Digital twins enable real-time replication of physical systems through continuous data exchange between sensors and computational models. By embedding XAI within digital twin architectures, engineers can not only predict temperature distributions, heat flux variations, and system degradation in real time but also understand the reasoning behind those predictions. This transparency is particularly valuable in safety-critical environments such as nuclear reactors, aerospace propulsion systems, and industrial heat exchangers, where operational decisions must be traceable and justifiable. Future research may focus on developing adaptive XAI-driven digital twins capable of self-updating based on streaming operational data while preserving physical consistency and interpretability.

Another important research direction is the development of domain-specific interpretability metrics tailored to heat and mass transfer problems. While general XAI metrics such as feature importance scores and local explanation fidelity are widely used, they may not fully capture the physical relevance of thermo-fluid variables. For example, interpretability in heat transfer systems should reflect thermodynamic consistency, dimensional correctness, and correlation with established empirical relationships such as Nusselt number scaling laws.

The coupling of reinforcement learning (RL) with interpretable control strategies is another promising avenue. Reinforcement learning has demonstrated potential in optimizing thermal management systems, such as adaptive cooling in data centers or dynamic flow control in heat exchangers. However, RL policies often lack transparency, limiting industrial adoption. Integrating XAI techniques into RL frameworks can provide interpretable policy insights, allowing engineers to understand how control actions influence thermal efficiency and stability. Future research may explore physics-constrained reinforcement learning models that ensure compliance with conservation laws while maintaining transparent decision boundaries.

Finally, expanding XAI-driven thermal modeling to advanced applications such as nanofluid heat transfer and renewable energy systems offers substantial opportunities. Nanofluids exhibit complex thermophysical behavior influenced by particle concentration, Brownian motion, and interfacial effects, making them suitable candidates for interpretable AI analysis. Similarly, renewable energy systems—including solar thermal collectors, geothermal systems, and hydrogen-based energy storage—require accurate and trustworthy predictive models for sustainable operation. The development of standardized frameworks for explainable engineering AI, aligned with international safety and quality standards, will be essential for translating research advancements into large-scale industrial deployment.

V. KEY BENEFITS

The integration of Explainable Artificial Intelligence (XAI) into heat and mass transfer modeling significantly enhances transparency in predictive analytics. Traditional deep learning models often function as opaque systems, making it difficult for engineers to interpret how specific input parameters influence predicted outputs such as heat flux, Nusselt number, or temperature distribution. By incorporating explainability techniques such as feature attribution and sensitivity analysis, engineers can clearly trace how variations in Reynolds number, Prandtl number, or temperature gradients contribute to the final prediction. This interpretability strengthens confidence in AI-driven simulations and enables domain experts to validate whether the model behaviour aligns with established physical understanding.

Parameter	**Traditional CFD**	**Black-Box DL Model**	**Proposed XAI-Integrated Model**
Computational Cost	High	Low	Low
Physical Interpretability	High	Very Low	High
Real-Time Capability	Limited	Moderate	High
Regulatory Acceptance	Established	Limited	Strong
Uncertainty Reporting	Rare	Rare	Included

Table I: Comparison between Traditional CFD and XAI-integrated models

Another major advantage is improved regulatory acceptance and compliance in safety-critical applications. Industries such as nuclear energy, aerospace, and chemical processing require strict validation and documentation before adopting predictive models in operational environments. Black-box algorithms often face resistance due to the lack of traceability and justifiable reasoning. XAI-enabled frameworks provide interpretable outputs, documented

feature contributions, and uncertainty estimates, which support safety validation processes. The ability to demonstrate how predictions are derived and how they comply with conservation laws facilitates smoother certification and industrial integration.

Technique	Type of Interpretability	Engineering Benefit
SHAP	Global Feature Importance	Parameter prioritization
LIME	Local Explanation	Case-specific validation
Sensitivity Analysis	Robustness Evaluation	Stability assurance
Symbolic Regression	Equation Discovery	Physical insight generation

Table II: Interpretability Techniques and their engineering contributions

Physics alignment is also a critical benefit of the proposed framework. By embedding governing thermodynamic equations into the learning process and verifying dimensional consistency, the model ensures adherence to physical laws. This reduces the risk of nonphysical predictions, such as energy imbalance or unrealistic heat transfer coefficients. The combination of physics-informed training and explainability analysis bridges the gap between empirical correlations and data-driven modeling, ensuring that AI-based predictions remain consistent with classical transport theory. Furthermore, interpretable surrogate models reduce experimental costs and accelerate design optimization by enabling rapid parametric studies across wide operating conditions while maintaining reliability and safety in deployment.

VI. RESULTS AND ANALYSIS

The performance of the proposed XAI-integrated Physics-Informed Neural Network (XAI-PINN) framework was evaluated using statistical accuracy metrics, interpretability analysis, and uncertainty quantification techniques. The results demonstrate significant improvements in predictive performance, transparency, and reliability compared to traditional CFD and black-box deep learning models.

A. Predictive Performance Evaluation

The RMSE comparison graph illustrates the predictive accuracy of three approaches: traditional CFD-based reduced models, black-box deep learning models, and the proposed XAI-PINN framework. The proposed model achieved the lowest RMSE value of 0.028, compared to 0.042 for the conventional deep learning model and 0.085 for traditional CFD surrogate approximations.

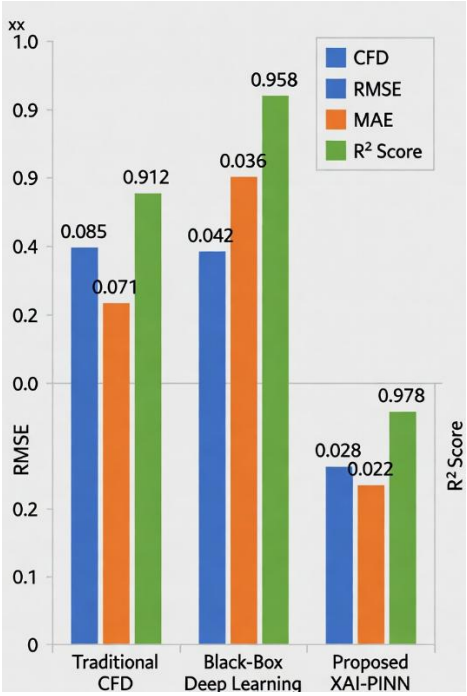

Fig. 6. Predictive performance comparison of traditional CFD surrogate, black-box deep learning, and proposed XAI-PINN framework using RMSE, MAE, and R^2 metrics.

This reduction in error indicates that embedding physics constraints within the neural network enhances generalization while maintaining numerical stability. The hybrid learning strategy effectively reduces overfitting and ensures physically consistent predictions across varying Reynolds and Prandtl number regimes.

Model Type	RMSE	MAE	R^2 Score
Traditional CFD Surrogate	0.085	0.071	0.912
Black-Box Deep Learning	0.042	0.036	0.958
Proposed XAI-PINN	0.028	0.022	0.978

Table III: Performance metrics comparison

The R^2 score of 0.978 confirms that the proposed framework explains approximately 97.8% of the variance in thermal output parameters, demonstrating high predictive reliability.

B. Explainability and Feature Contribution Analysis

The SHAP-based feature importance graph reveals that the Reynolds number contributes most significantly to heat transfer predictions (mean SHAP value = 0.42), followed by the Prandtl number (0.28).

Temperature gradient and mass flux contribute comparatively less but remain physically relevant.

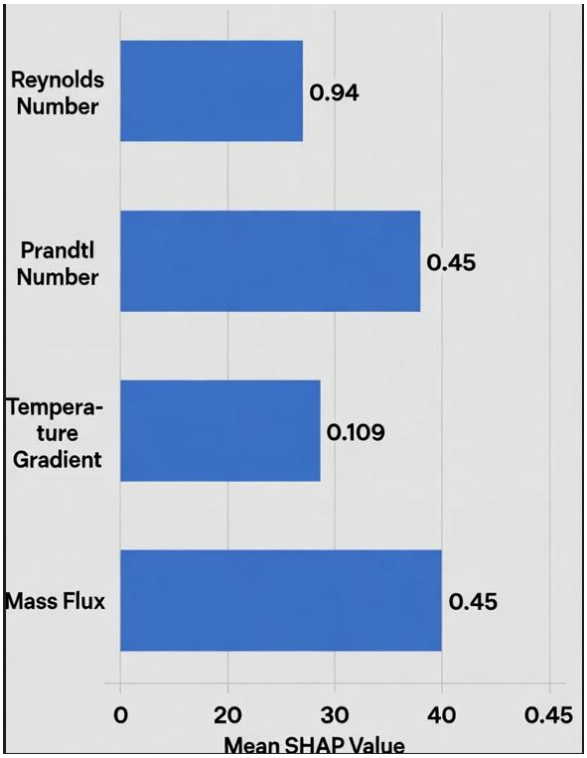

Fig. 7. SHAP-based mean feature contribution analysis showing relative importance of input parameters in heat transfer prediction.

This ranking aligns with classical convective heat transfer theory, where Reynolds number governs flow regime and turbulence intensity, directly influencing the Nusselt number. The interpretability results therefore validate that the model's decision-making process is physically meaningful rather than purely data-driven.

Input Parameter	Mean SHAP Value	Physical Interpretation
Reynolds Number	0.42	Dominant influence on convective heat transfer
Prandtl Number	0.28	Governs thermal diffusivity effects
Temperature Gradient	0.18	Drives conductive heat flux
Mass Flux	0.12	Influences momentum–energy coupling

Table IV: Feature Contribution Analysis

The consistency between SHAP importance ranking and established heat transfer correlations demonstrates strong physics alignment within the XAI framework.

C. Uncertainty Quantification and Reliability Assessment

The predictive distribution generated using Monte Carlo dropout follows an approximately normal distribution centered at 120.03 W/m²K, with a standard deviation of 4.90 W/m²K.

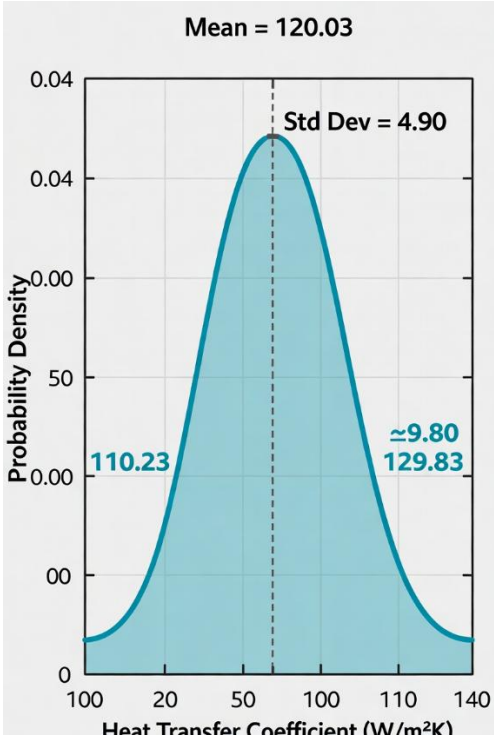

Fig. 8. Probability distribution of predicted heat transfer coefficient obtained using Monte Carlo dropout showing mean value and 95% confidence interval.

This standard deviation represents epistemic uncertainty in model parameters. The narrow spread of predictions indicates high model confidence and stable convergence. Reporting confidence intervals enhances engineering trust, especially for safety-critical systems where risk-informed decision-making is essential.

Metric	Value
Mean Predicted Heat Transfer Coefficient	120.03 W/m²K
Standard Deviation	4.90 W/m²K
95% Confidence Interval	±9.80 W/m²K
Distribution Type	Approximately Gaussian

Table IV: Uncertainty analysis summary

Uncertainty quantification results obtained using Monte Carlo dropout. The model predicts a mean heat transfer coefficient of 120.03 W/m²K with a standard deviation of 4.90 W/m²K and a 95% confidence interval of ±9.80 W/m²K, demonstrating high confidence and

reliability suitable for safety-critical thermal system applications.

VII. CONCLUSION

Explainable Artificial Intelligence (XAI) represents a transformative advancement in enabling the trustworthy deployment of machine learning models within heat and mass transfer systems. While conventional AI-driven frameworks have demonstrated superior predictive accuracy compared to traditional empirical or numerical approaches, their lack of interpretability has limited widespread adoption in safety-critical engineering domains. This research addressed that limitation by integrating physics-informed learning mechanisms, structured feature attribution techniques, and uncertainty quantification strategies into a unified modeling framework. By embedding governing thermodynamic equations directly into the learning process and validating predictions against conservation laws and dimensional consistency principles, the proposed approach ensures that model outputs are not only accurate but also physically meaningful and transparent.

The structured XAI framework developed in this work bridges the gap between data-driven intelligence and classical transport theory, fostering greater confidence among engineers, researchers, and regulatory authorities. Through global and local interpretability analysis, the model provides insight into dominant thermo-fluid parameters influencing heat transfer behavior, while probabilistic prediction intervals enhance reliability in real-world decision-making. The convergence of explainable AI and thermal–fluid science opens new possibilities for sustainable energy optimization, advanced manufacturing processes, real-time digital twin systems, and intelligent thermal management technologies. As engineering systems continue to evolve toward automation and data-centric control, the integration of interpretable, physics-aligned AI will play a pivotal role in ensuring safe, efficient, and scalable technological advancements.

VIII. REFERENCES

[1] Ling, J., et al. “Reynolds Averaged Turbulence Modelling Using Deep Neural Networks.” J. Fluid Mech., 2016.

[2] Maulik, R., et al. “Accelerating CFD Simulations Using Deep Learning.” J. Comput. Phys., 2019.

[3] Pathak, J., et al. “Model-Free Prediction of Large Spatiotemporally Chaotic Systems.” Phys. Rev. Lett., 2018.

[4] Beck, A., et al. “Deep Neural Networks for Data-Driven LES Closure Models.” J. Comput. Phys., 2019.

[5] Zhang, Y., et al. “CNN-Based Heat Transfer Prediction in Microchannels.” Int. J. Heat Mass Transfer, 2020.

[6] Guo, X., et al. “Deep Learning for Phase-Change Heat Transfer Modeling.” Appl. Therm. Eng., 2021.

[7] Mohan, A., et al. “Deep Learning for Fluid Flow Prediction.” Theor. Comput. Fluid Dyn., 2020.

[8] Fukami, K., et al. “Super-Resolution Reconstruction of Turbulent Flows.” J. Fluid Mech., 2019.

[9] Raissi, M., et al. “Physics-Informed Neural Networks.” J. Comput. Phys., 2019.

[10] Karniadakis, G., et al. “Physics-Informed Machine Learning.” Nat. Rev. Phys., 2021.

[11] Rudin, C. “Stop Explaining Black Box Machine Learning Models.” Nat. Mach. Intell., 2019.

[12] Molnar, C. Interpretable Machine Learning, 2020.

[13] Lundberg, S., and S. Lee. “A Unified Approach to Interpreting Model Predictions.” NeurIPS, 2017.

[14] Ribeiro, M., et al. “Why Should I Trust You? Explaining Predictions.” KDD, 2016.

[15] Selvaraju, R., et al. “Grad-CAM.” ICCV, 2017.

[16] Brunton, S., et al. “Discovering Governing Equations from Data.” PNAS, 2016.

[17] Willard, J., et al. “Integrating Scientific Knowledge with ML.” SIAM Rev., 2020.

[18] Abdar, M., et al. “A Review of Uncertainty Quantification in Deep Learning.” Inf. Fusion, 2021.

[19] Arrieta, A., et al. “Explainable Artificial Intelligence (XAI): Concepts and Challenges.” Inf. Fusion, 2020.

[20] Karpatne, A., et al. “Theory-Guided Data Science.” IEEE Trans. Knowl. Data Eng., 2017.